



# EVENT SOURCING WITH ELASTICSEARCH

“ES with ES”

by Paweł Świątkowski

# Agenda


- Event Sourcing – what and why?
- ElasticSearch – what and why?
- Putting it together - how we do it @ Boostcom

# About me

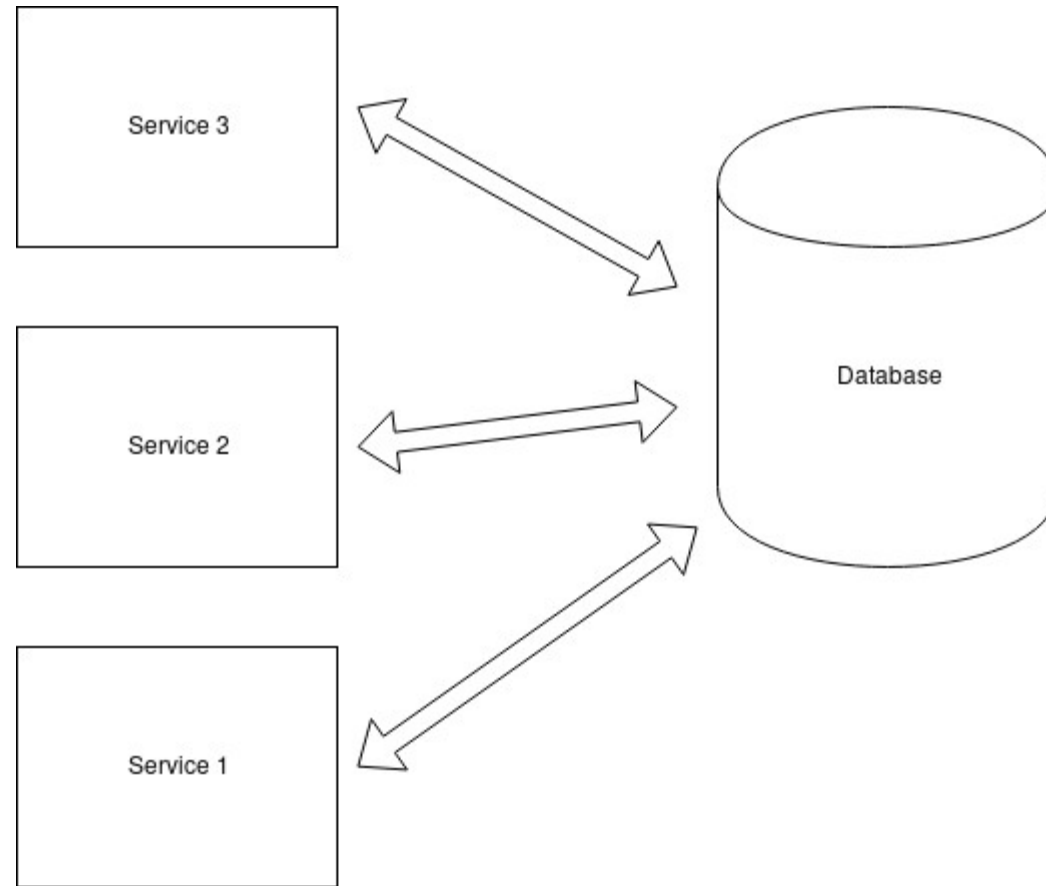
- Ruby programmer at Boostcom
- Blog: <http://katafrakt.me>
- Rails-free Ruby newsletter: <https://rubytuesday.katafrakt.me>
- NOT an Event Sourcing expert or theorist

# EVENT SOURCING

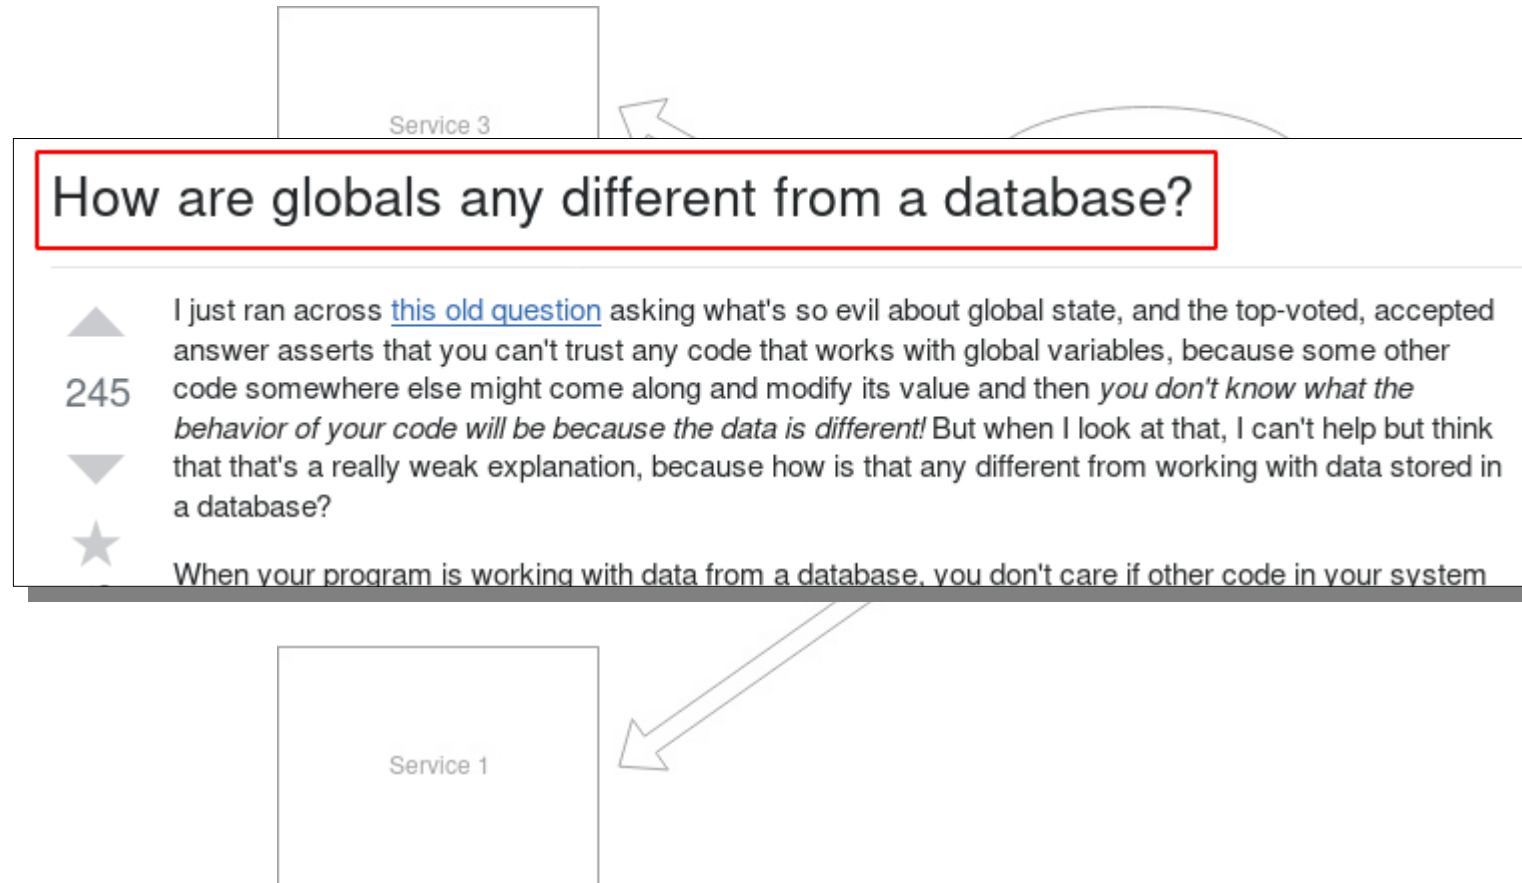
# Typical web application



# Typical modern web application



# Typical modern web application



# Typical modern web application as code

```
$state = [{ id: 1, name: 'Element 1' }, { id: 2, name: 'Element 2' }]

def mutate!(i)
  $state.shuffle!
  $state[0][:name] = "Element #{i}"
end

5.times.map do |i|
  Thread.new { mutate!(i) }
end.each(&:join)
```



# HOW DID WE END UP HERE?!?!?

- Logs
- Papertrail
- Audit log



What if we could have  
temporal data (history of  
changes) as first-class data?

# Events

- { event: 'name\_change', name: 'Element 2', id: 1, timestamp: 12345 }
- { event: 'name\_change', name: 'Element 1', id: 1, timestamp: 12346 }
- { event: 'name\_change', name: 'Element 5', id: 2, timestamp: 12347 }
- { event: 'name\_change', name: 'Element 2', id: 1, timestamp: 12348 }
- { event: 'shuffle\_array', ids: [2, 1], timestamp: 12349 }

# Events Sourcing rules

- Only append
- Always have timestamp
- Event source / event log is your single source of truth

# Events Sourcing wins

- You always have the history
- History is first-class data, you can perform aggregations etc.
- In rare cases you can get system to the state of some time in the past and debug
- You always know where to seek the truth


# Events Sourcing losses

- Harder to get current state
- Take up much more storage space
- Different mindset – harder to get people on board

# Traditional tools for event store



***EVENT STORE***®



***cassandra***



mongo**DB**



elasticsearch



# ElasticSearch – what it is?

- Built in 2010 as full-text search engine
- Widely used as E in ELK Stack
  - Elasticsearch – Logstash – Kibana
  - Centralized logging
- Basically a document store

# ElasticSearch

## Pros

- Can easily handle large amount of data
- Optimized for logs – append-only structure
- Good search and filtering capabilities
- Scripting in Groovy
- Kibana for free

## Cons

- Complicated to set up properly
- Terrible query language
- Extensive scripting can kill the server

Average conference talk about Event Sourcing, Q&A part

Q: How to begin with ES?

A: Oh, just start gathering  
some events...

2 years later...

OK, now what?

# HOW WE DO IT AT BOOSTCOM

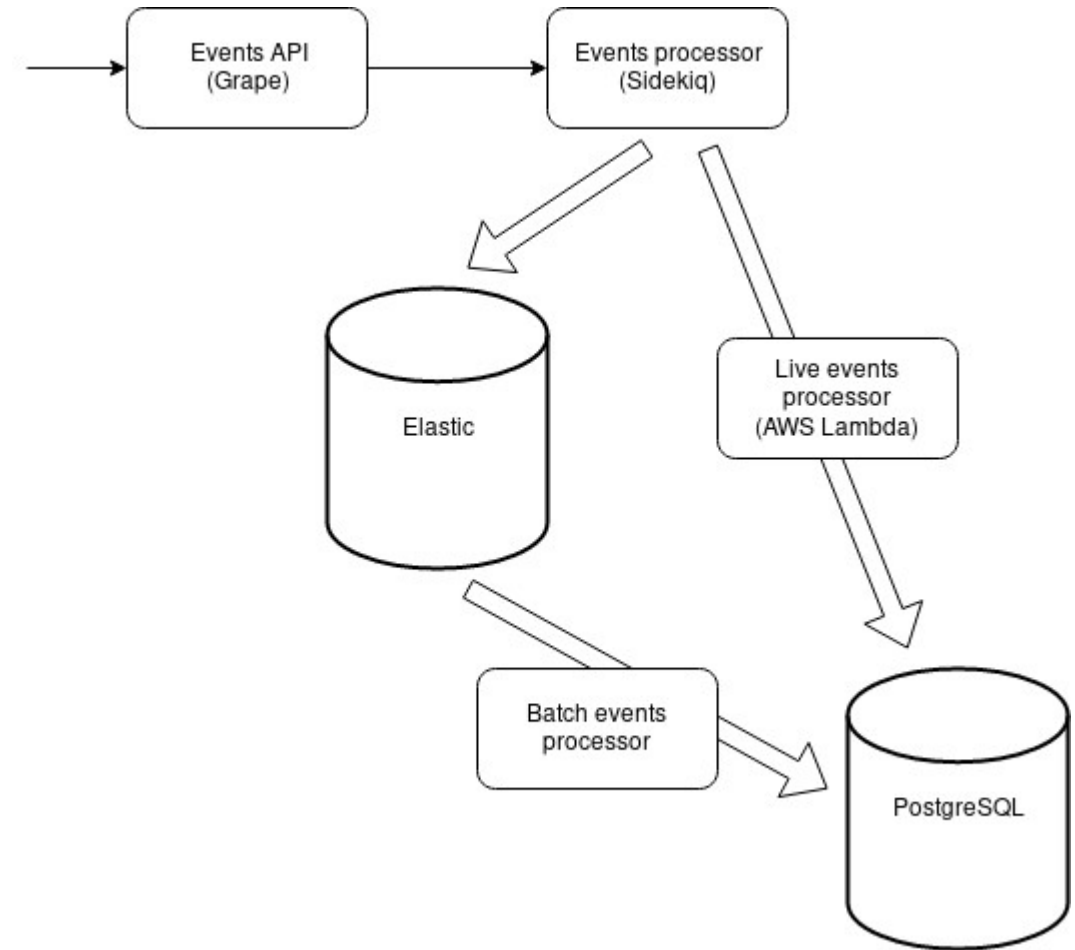
Putting it all together

# Some facts

- Our traffic is mostly peak-based
- We handle 2 – 3 million daily (**2448995** – exact daily average from last two weeks)
- Some external integrations are really bad at resending events when something happens
- Others send duplicates on regular basis

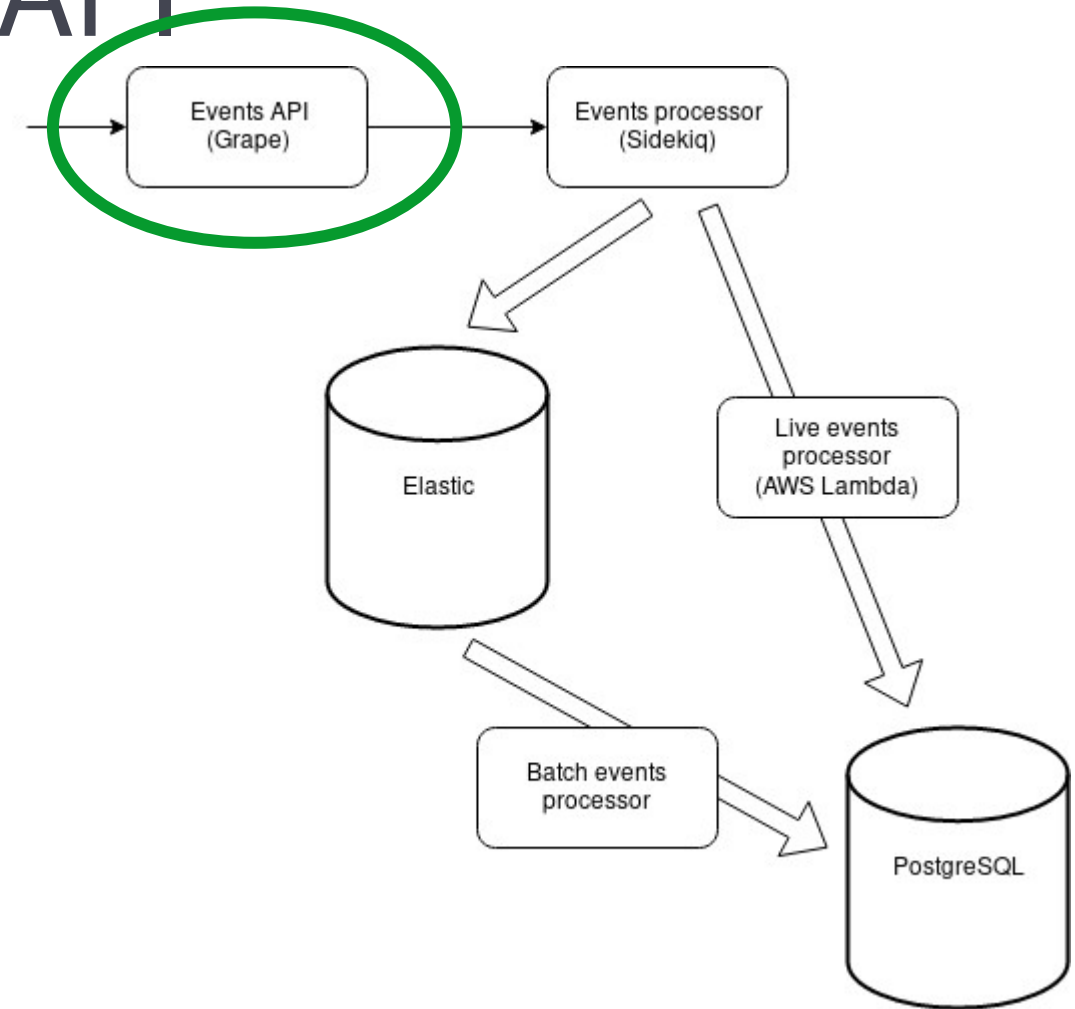
# Architecture

- 4 components
- 2 databases
- 1 web application
- 1 serverless set of functions




# Architecture – Events API

- Only web-facing app
- HTTP API on Grape
- Only authorization and brief validation
  - If event happened, it happened
- Saves Sidekiq jobs to be picked up asynchronously
- Never failed ;)





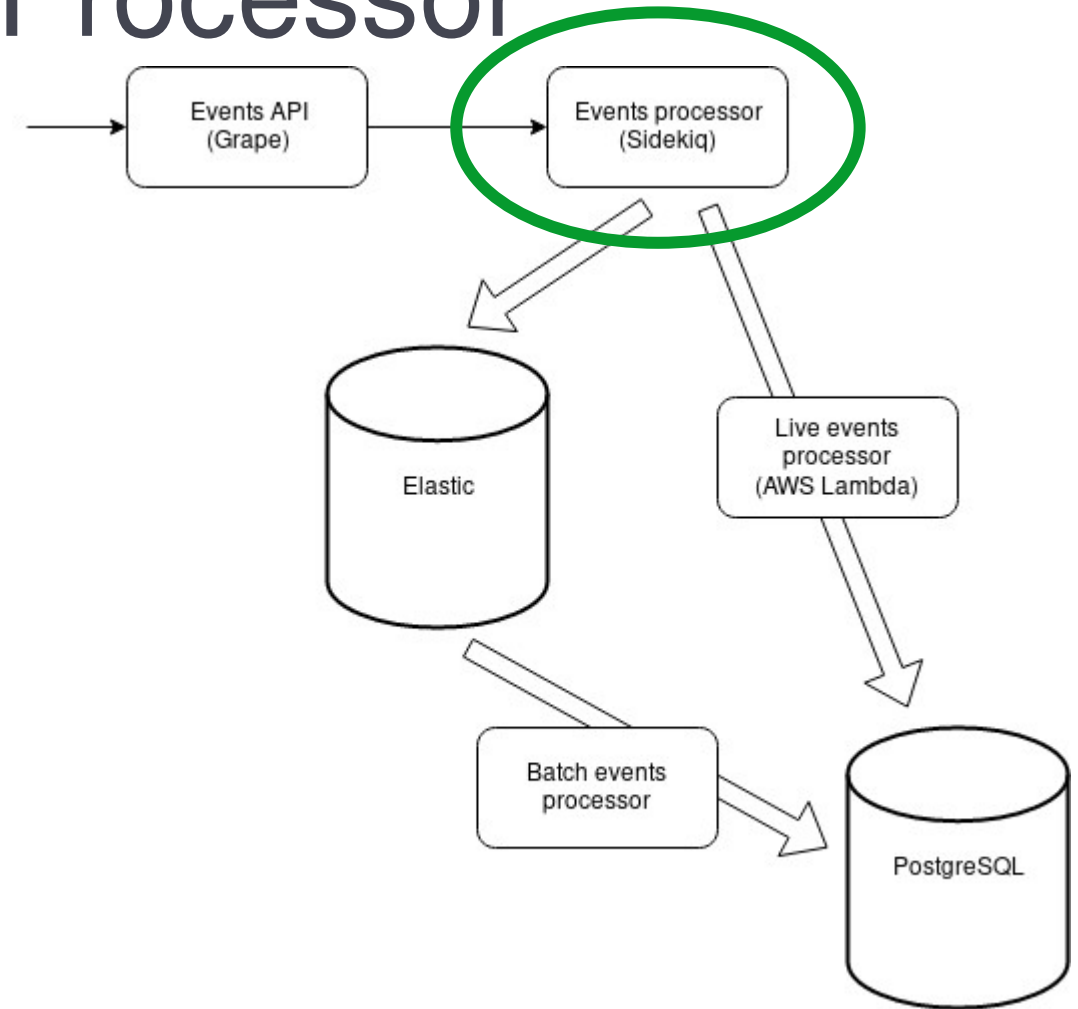
# Architecture – Events API



PostgreSQL

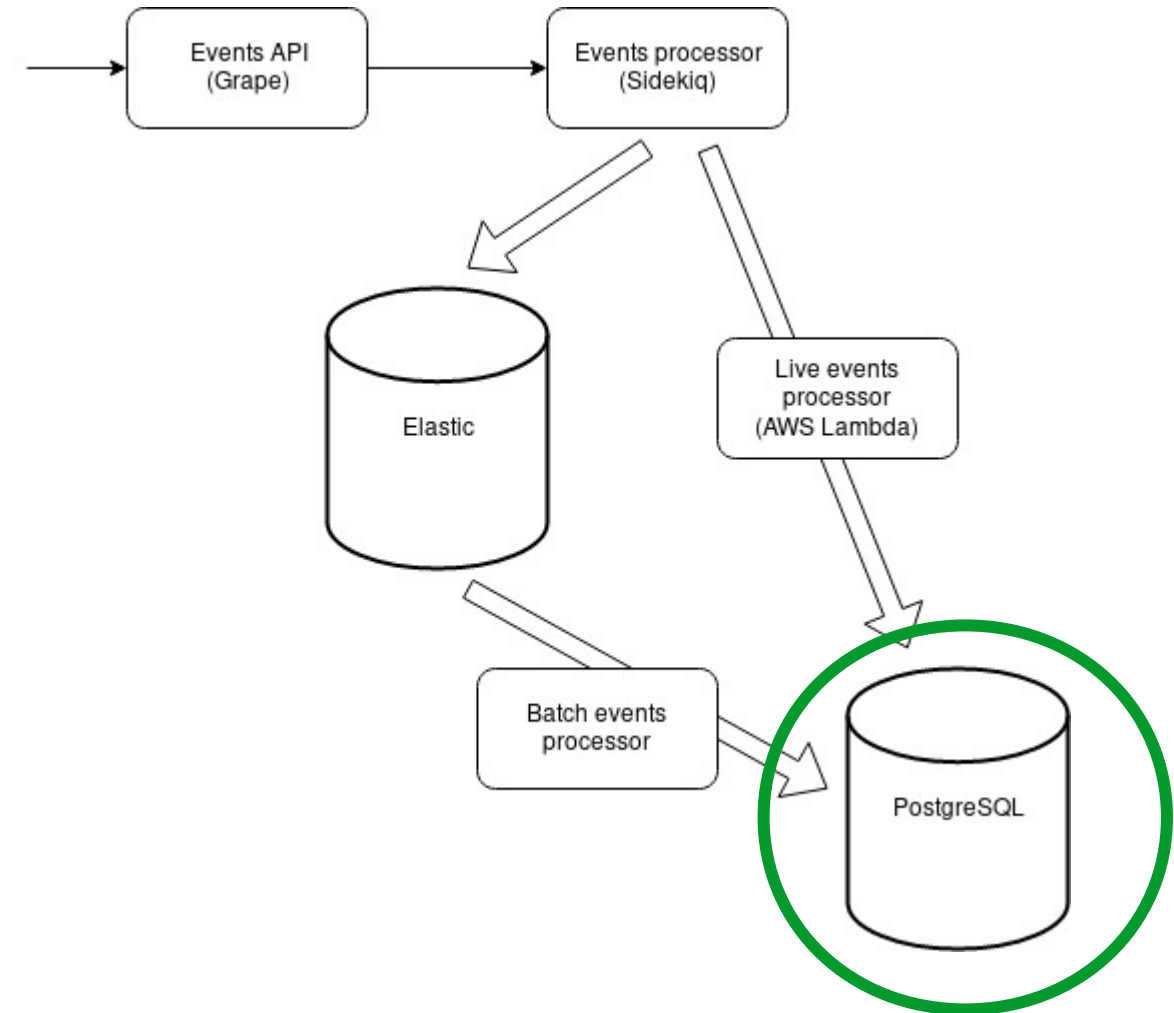
# Architecture – Events Processor

- Does all the heavy lifting
- Checks for duplicates (checksums)
- Adds missing data
- Creates internal events
- Saves to Elasticsearch and Kinesis



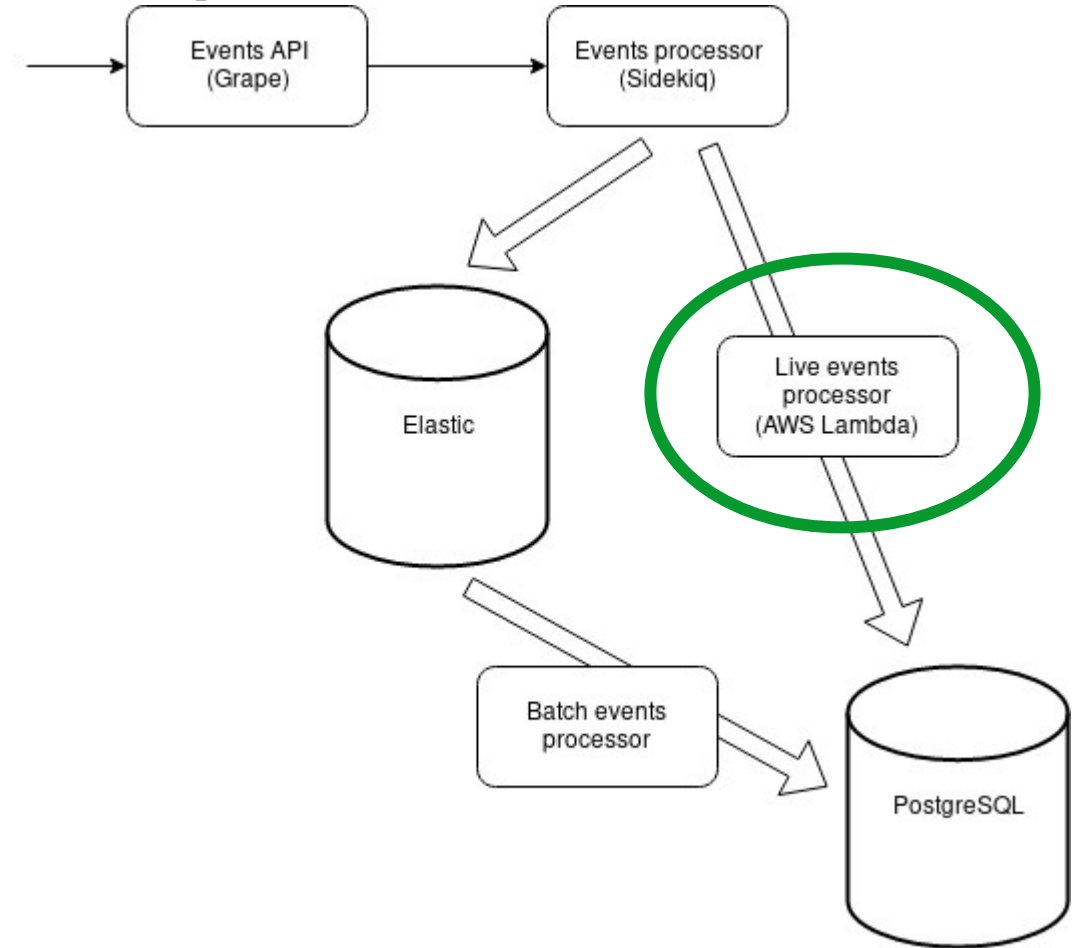
# Architecture – read database

- Events log is not very well-suited for complex querying (no joins)
- We store data in read-optimized SQL database too



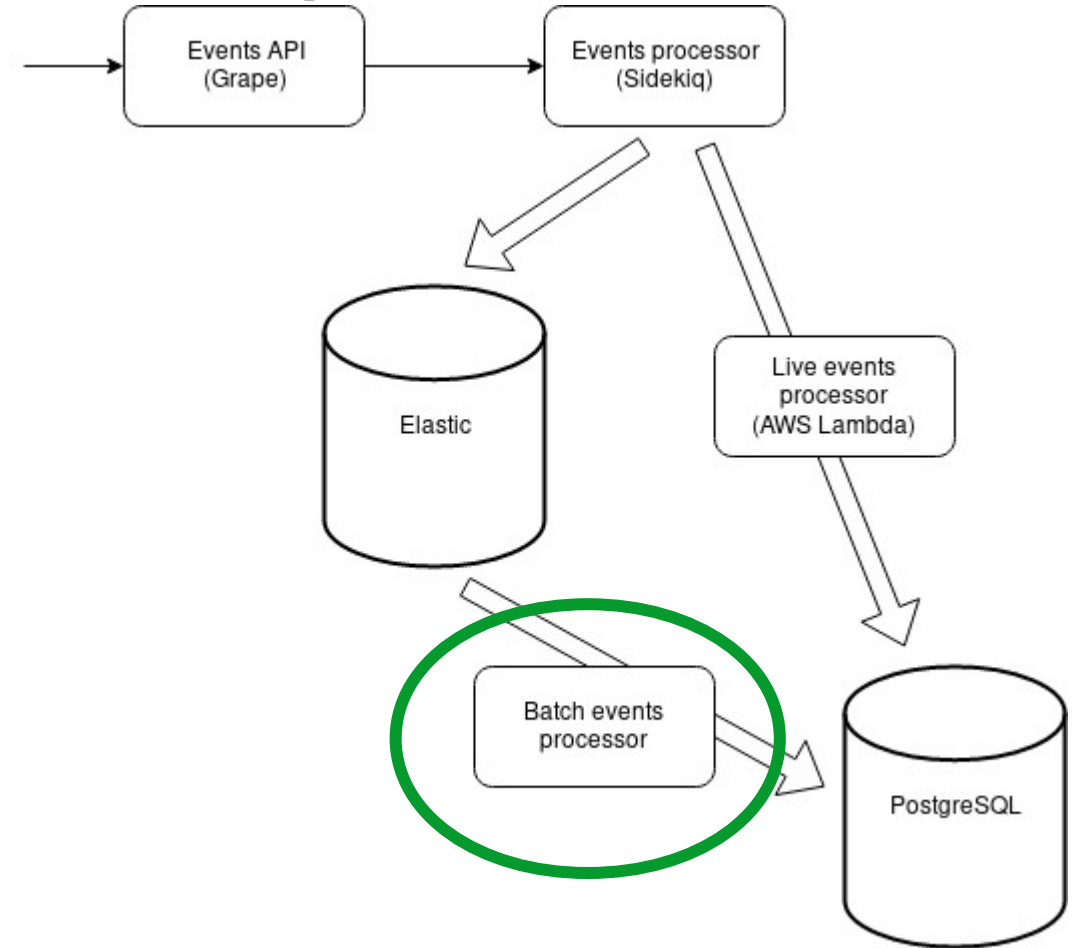
# Architecture – live events processor

- Takes an event and “applies it” to read database (for example mark that user used at least one coupon)
- Written in Python :(



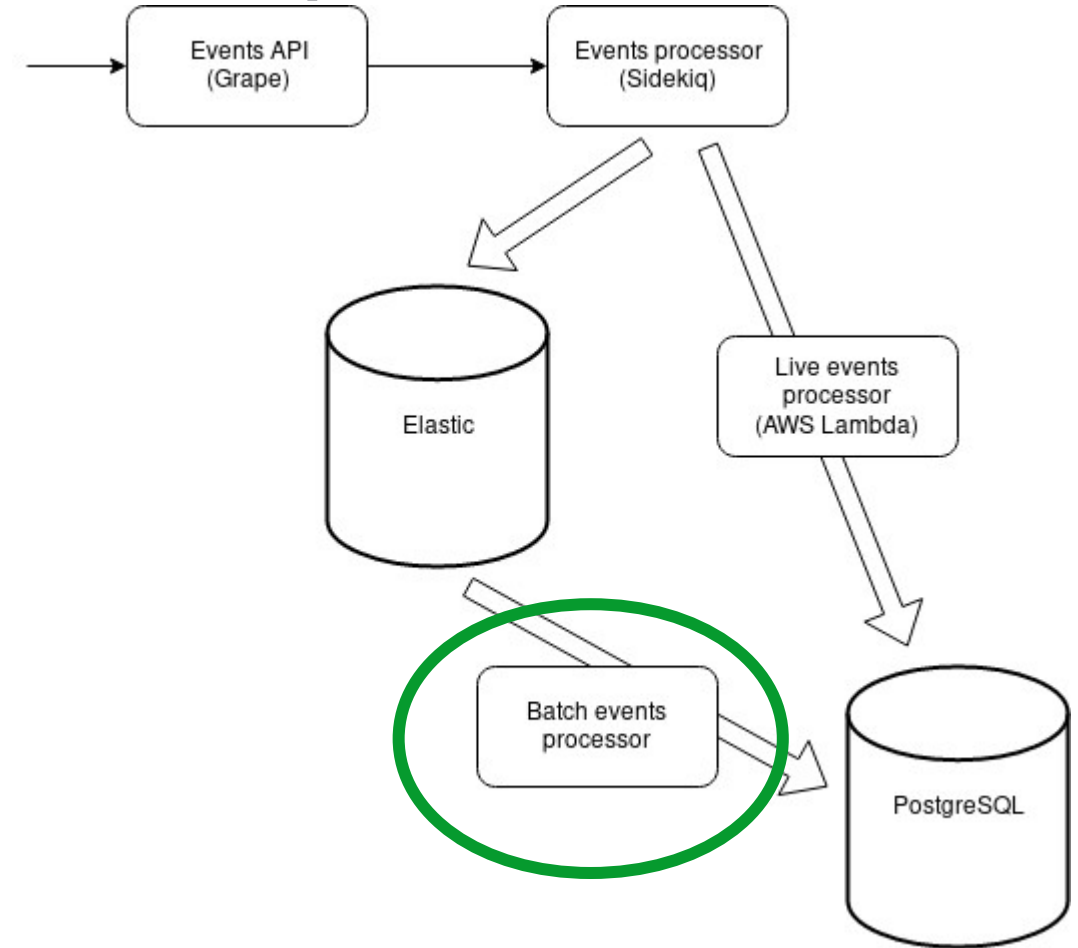
# Architecture – batch events processor

- Drop read database and rebuild from scratch!
- Is not just applying events in order (would be too slow)
- I also relatively slow
- Some code duplicated from live processor



# Architecture – batch events processor

- Written in Ruby
- Uses RocksDB for storing intermediate results
- Outputs CSVs which are COPY'd to PostgreSQL
- Uses Elasticsearch scrolls

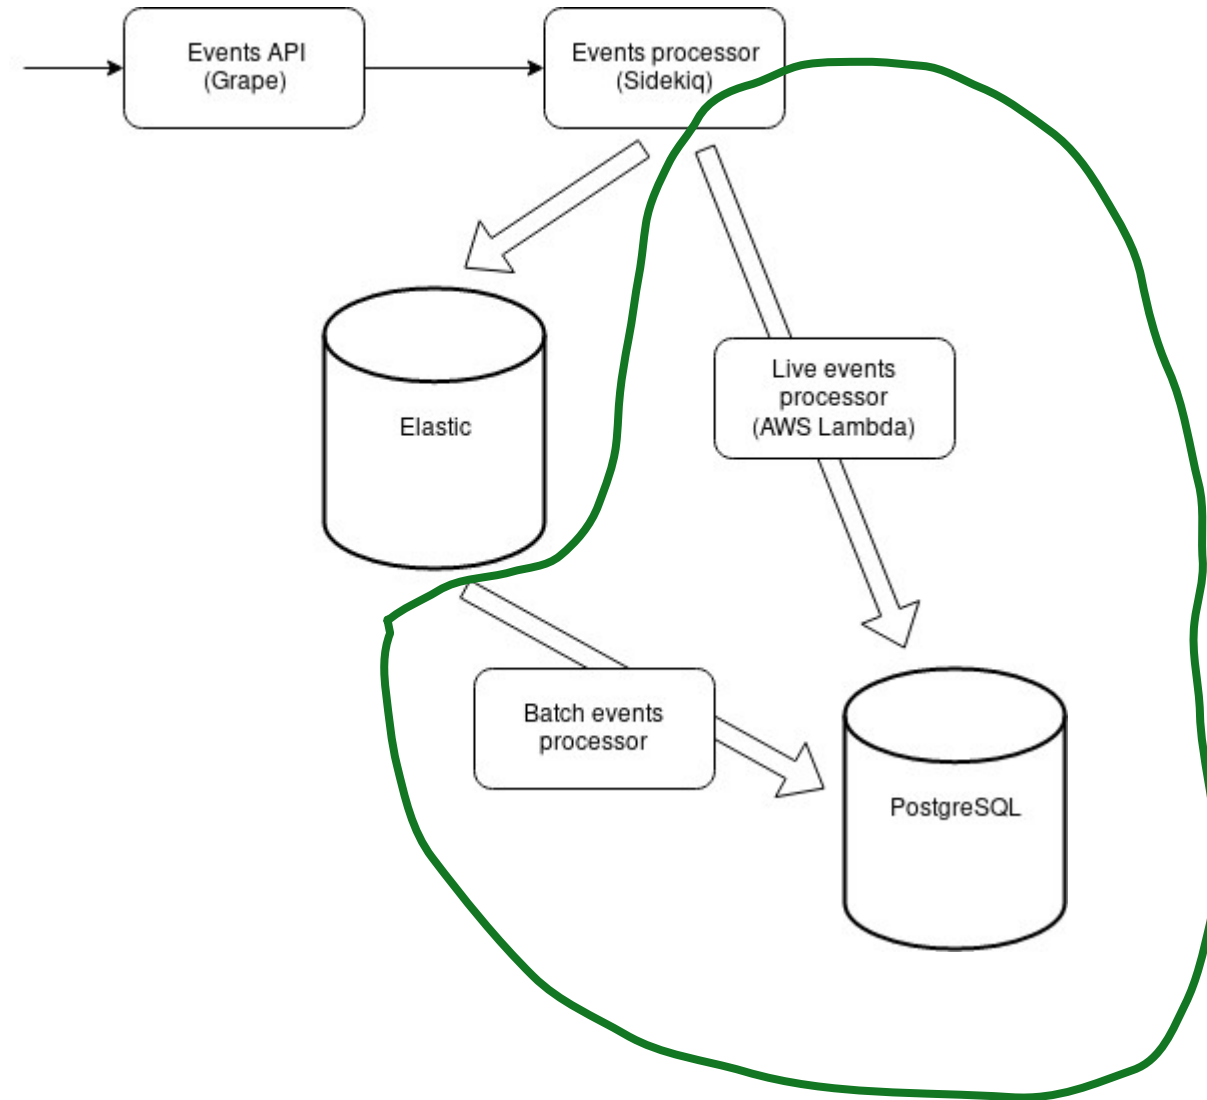


# Part where you can make mistakes

- Wrong database layout or missing column?
- Some events processed in wrong order?
- Database or Lambda down?

NO PROBLEM\*

(\* You're not losing data)

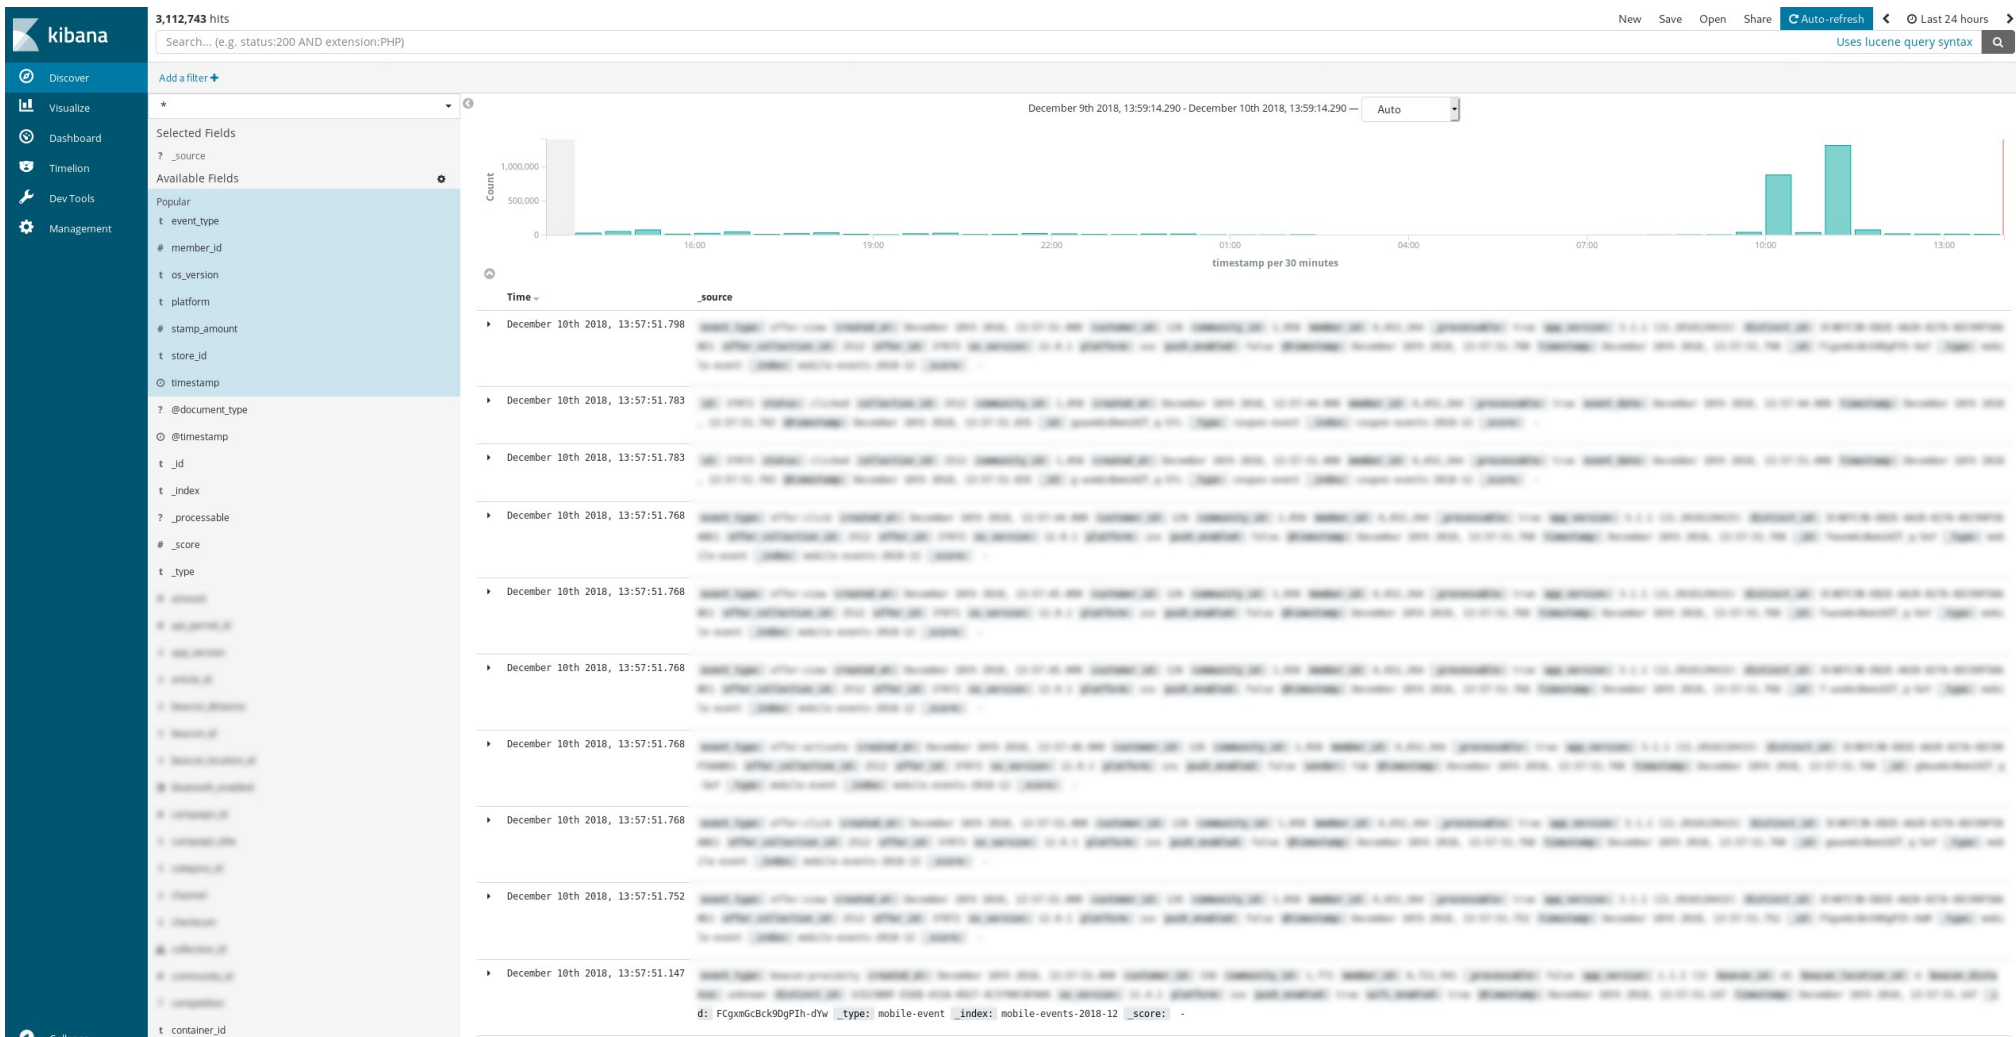


# Benefits of Elasticsearch


- Events in store are append only and immutable, but...
  - ES allows updates without rewriting the whole log (can be done live)
- We are not afraid about growing data volume
- Can connect to many existing solutions (PowerBI)
- Basic insights using Kibana



# Kibana – event filtering, viewing etc.



# Kibana – visualizations



# Kibana benefits

- We were able to teach some non-programmers to use Kibana before asking questions
- It's faster for us too to use it's interface for checking things
- You cannot break anything



# Tips & Tricks

- Don't try to event source everything. Because:
  - tech constraints (synchronous things, uniqueness checks...)
  - not real events
  - law (GDPR *et consortes*) - remember about personal data!
- Collect many events, not only ones you think you need now
- Don't be afraid to give up on some events if their quality is bad
- Consider seeding



WE'RE HIRING

¿QUESTIONS?